



# Introduction to Hive

Tuhin Mahmud  
@Big Data Revealed Meetup group  
10/11/2014

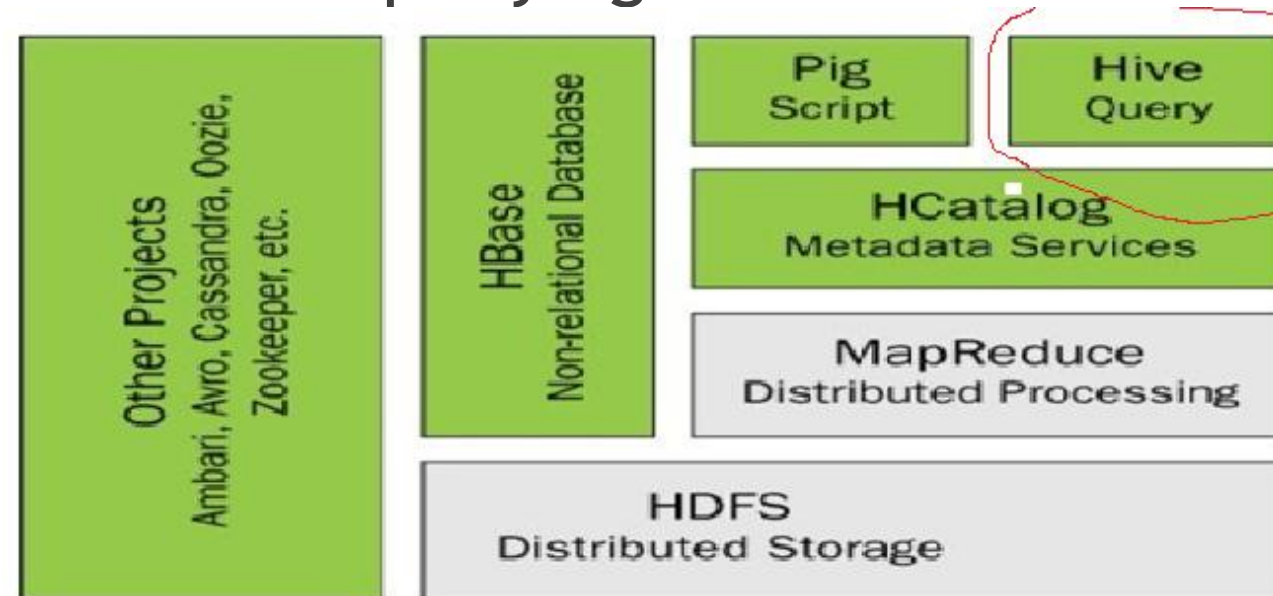
# About me

- ▶ Name: Tuhin Mahmud
- ▶ Software developer @IBM
- ▶ Cloudera Certified Hadoop Developer
- ▶ Email: [tuhinm@hotmail.com](mailto:tuhinm@hotmail.com)
- ▶ [www.linkedin.com/in/tuhinmahmud](http://www.linkedin.com/in/tuhinmahmud)

# What is Hive?



- Hive is an essential Hadoop ecosystem tool that provides an SQL dialect for querying data stored in Hadoop.[1]



The Hadoop 1.0 ecosystem.



# What can Hive do?

- ▶ Hive makes it easier to port **SQL** based application to Hadoop.
- ▶ Hive is most suited for **data warehouse** applications where relatively *static* data is analyzed.
- ▶ It **does not provide OLTP**( online transaction processing)
- ▶ Hive with Spark => **Shark** => **Spark SQL**

# Hive vs RDMS

## Hive

For **analytics** and **large** aggregation

Data WareHouse

High Latency, Fast load and flexibility

No transaction Support

## RDMS

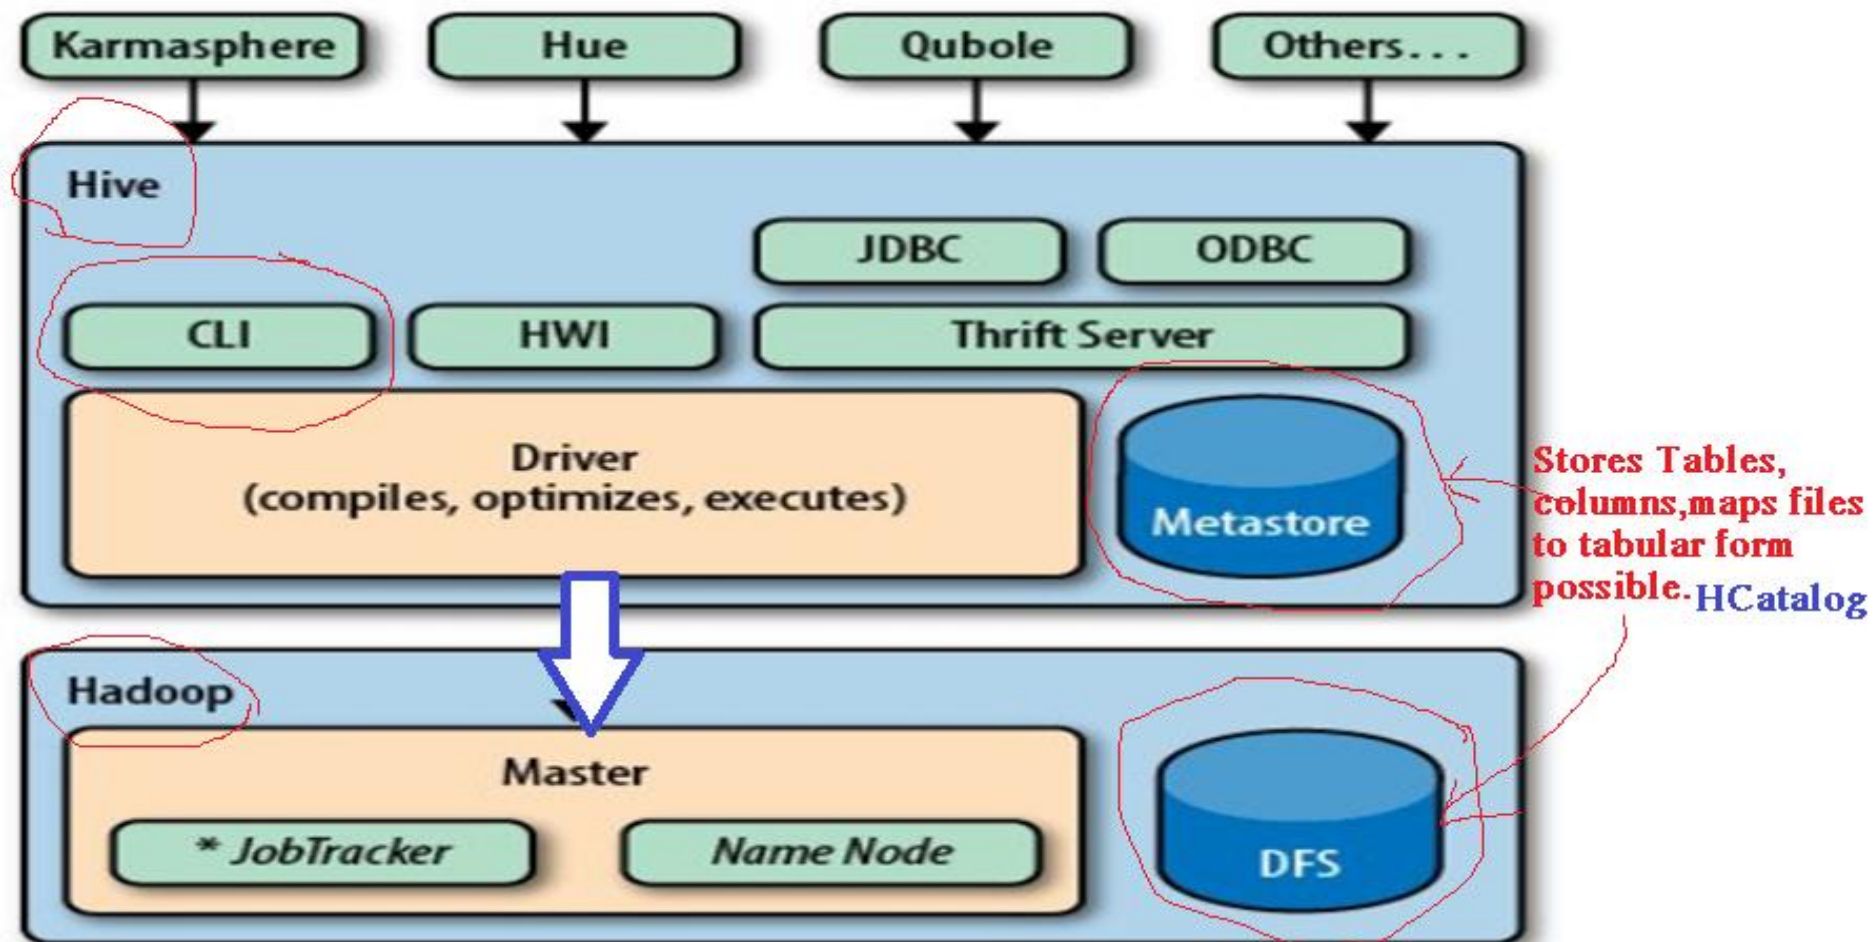
**Real time** processing

Database

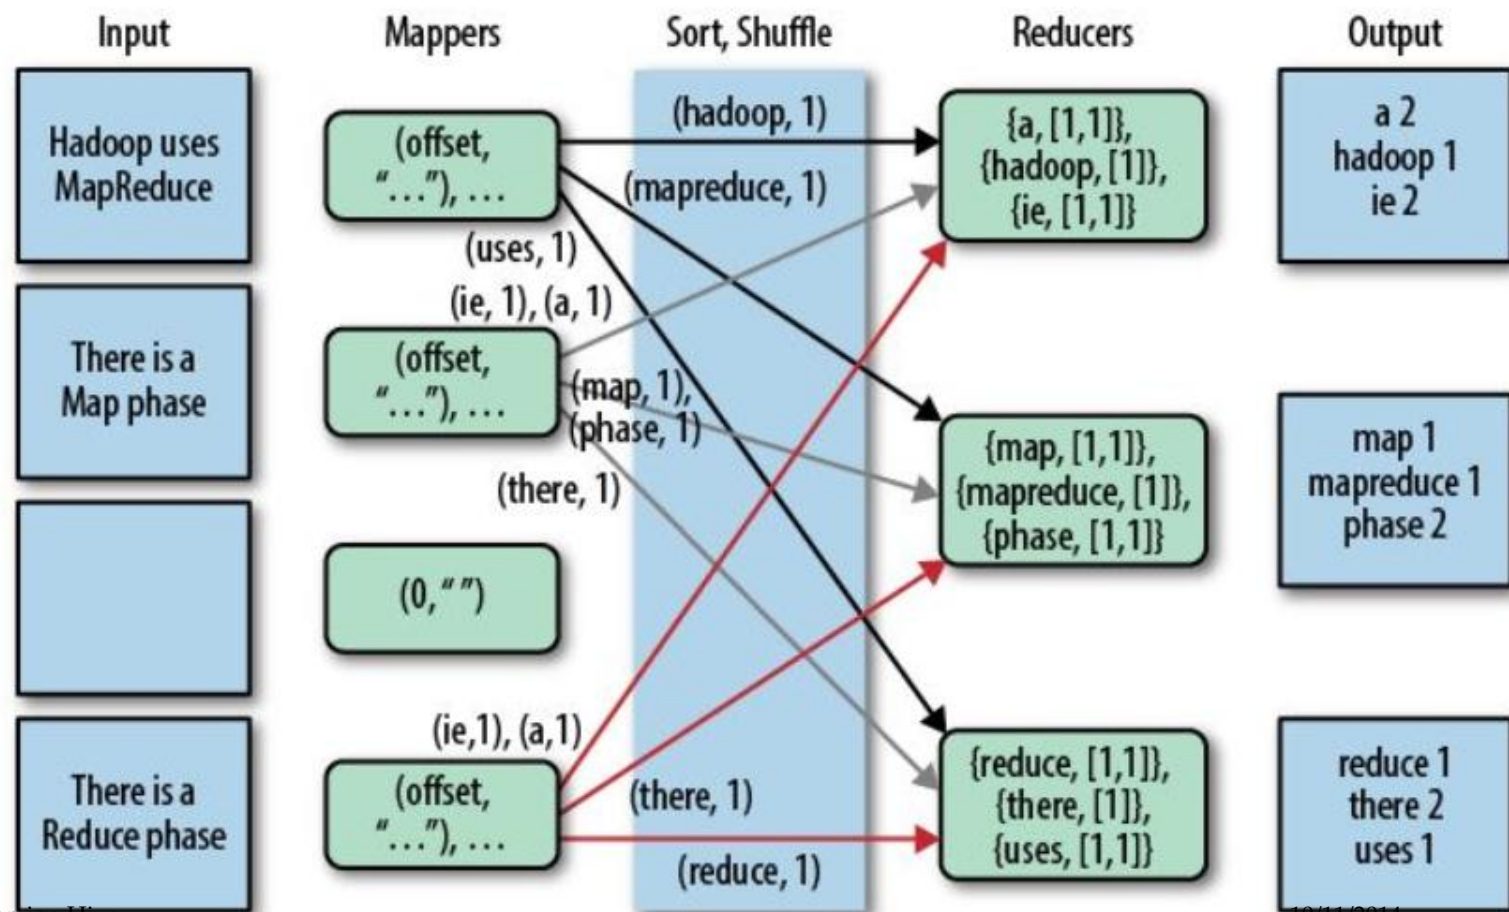
**Fast Query**

**Transaction** support

# Hive Components



# Word Count Example Revisited





## Word Count Example Code

- ▶ 63 lines of java code
- ▶ 25 lines of python code
- ▶ How many lines do you think we will need to write the same in Hive programming Language?



# Word Count Example Code

- ▶ 63 lines of java code
- ▶ 25 lines of python code
- ▶ How many lines do you think we will need to write the same in Hive programming Language?
- ▶ 3 lines and 1 line with pre-processing!

# Hive Directory Structure

## ▶ Lib Directory

- ▶ `$HIVE_HOME/lib`
- ▶ Location of Hive **JAR file**
- ▶ Contains the actual Java code that implement the Hive functionality

## ▶ Bin Directory

- ▶ `$HIVE_HOME/bin`
- ▶ Location of Hive Scripts/Services

## ▶ Conf directory

- ▶ `HIVE_HOME/conf`
- ▶ Location of configuration files

# Hive Metastore

- ▶ What is Hive Metastore?
  - ▶ Hive metastore is a database that stores **metadata** about the hive tables. Metadata examples are *table name*, *column name*, *column type*, *table location* etc.
- ▶ Datastore
  - ▶ In process ( **Deby**)
  - ▶ Out of process Datastores e.g DB2, MySQL, Oracle etc.
- ▶ Configuration for metastore:
  1. **Embedded**( same JVM)
  2. **Local** ( out of process database)
  3. **Remote** ( hive service JVM is separate from metastore JVM and database is remote place.

# Hive Data units

- ▶ Database
- ▶ Table
- ▶ Partition
  - ▶ fundamentally horizontal slices of data which allow large sets of data to be segmented into more manageable chunks.
- ▶ Buckets(Clusters)

# Partition example

```
CREATE TABLE customer (  
  id      INT, name      STRING, address1  STRING, address2  STRING,  
  city    STRING, state  STRING, zip      STRING  
)  
  
PARTITION BY (  
  REGION  STRING,  
  country STRING  
)  
;
```

## Directory Structure in hdfs ( for faster query)

/erp.db/customer/region=North America/country=US

/erp.db/customer/region=North America/country=CA

/erp.db/customer/region=South America/country=BR

# Cluster and Bucket example

```
CREATE TABLE order (  
  username  STRING,  
  orderdate STRING,  
  amount    DOUBLE,  
  tax       DOUBLE,  
) PARTITIONED BY (company STRING)  
CLUSTERED BY (username) INTO 25 BUCKETS;
```

- We are creating 25 **buckets** and **clustering** on 'username'

# Physical Layout

- ▶ Data files are regular HDFS files
- ▶ Warehouse Directory in HDFS Specified in hive-site.xml as `hive.metastore.warehouse.dir`  
e.g `/home/hive/warehouse`
- ▶ Tables stored in **subdirectories** of the warehouse
- ▶ **Partion** and **buckets** subdirectories of Table subdirectory.
- ▶ Data stored in flat file in **HDFS** with char delimited text or sequence file.



# Hive DDL commands

## ► Create database

```
CREATE DATABASE mydatabase;
```

```
CREATE DATABASE mydatabase LOCATION '/myfolder/subdir/;
```

```
CREATE DATABASE mydatabase COMMENT 'This is my database';
```

## ► Delete database

```
DROP DATABASE IF EXISTS mydatabase;
```

# Data Types

- ▶ **Primitive types**

- Integers**

- Boolean**

- String**

- Date/Time**

- Binary**

- ▶ **Complex Types**

- Arrays**

- Structs**

- Maps**

- Union**

# Hive Operators

## ▶ Relational Operators

- ▶ `SELECT id,name FROM users WHERE name LIKE 'Tom%';`
  - ▶ Return values Tom ,Tomas

## ▶ Logical Operators

- ▶ `AND ,&&, OR, ||, NOT A , !A`
- ▶ Operators on complex Types
- ▶ `A[n]` for array access
- ▶ `M[key]` for map access
- ▶ `S.x` for struct field

# Hive functions

## ► Built in Functions

- count(\*)
- sum(col)
- avg(col)
- min(col),max(col)

```
SELECT sum(price) FROM fruits;  
>27.75
```

```
SELEECT count(*) FROM furits;  
> 4
```

Comprehensive list of available functions

<https://cwiki.apache.org/confluence/display/Hive/LanguageManual+UDF>

Table “**fruits**”

apple	5.25
oranges	10.00
pipeapple	5.00
grapes	7.50

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern, layered effect on the right side of the slide.

Code & demo

# Hive Command line

## ► CLI ( Command line interface)

*>hive*

*-d,--define <key=value> Variable substitution to apply to hive commands. e.g. -d A=B or --define A=B*

*-e <quoted-query-string> SQL from command line*

*-f <filename> SQL from files*

*-H,--help Print help information*

## ► Beeline - New Command Line Shell

<https://cwiki.apache.org/confluence/display/Hive/HiveServer2+Clients#HiveServer2Clients-Beeline-NewCommandLineShell>

# Hive Clients

- ▶ Command Line
- ▶ JDBC
- ▶ Python
- ▶ PHP
- ▶ ODBC

<https://cwiki.apache.org/confluence/display/Hive/HiveClient>



# Hive FAQs

- ▶ What is the default size of data processed per reducer in Hive?
  - ▶ **1 G** controlled by using *hive.exec.reducers.bytes.per.reducer*
- ▶ How can I control number of mappers and reducers in Hive?
  - ▶ *mapred.max.split.size* and *mapred.max.split.size*.
  - ▶ *hive.exec.reducers.bytes.per.reducer*.
- ▶ *Check the reference below for more similar FAQs*
  - ▶ <https://developer.ibm.com/hadoop/docs/getting-started/faqs/#hivfaq>

# Hive Alternatives

- ▶ **HBASE** - for rapid query and row level update ,transaction support
- ▶ **Pig** -High level data flow language
  - ▶ Higher learning curve for SQL programmers
  - ▶ Good for ETL , not good for ad-hoc query
  - ▶ Used in combination with Hive
- ▶ **Java Map Reduce**
  - ▶ Java code for simple word count is about 63 lines.
  - ▶ More Hadoop internal architecture know how needed and quick prototyping difficult.

## Other choices

- ▶ BigSQL ,BigR- from IBM
- ▶ Impala - from [www.cloudera.com](http://www.cloudera.com)
- ▶ MapR- from [www.mapr.com](http://www.mapr.com)
- ▶ Cassandra - [cassandra.apache.org](http://cassandra.apache.org)

# Big SQL

- ▶ What is Big SQL?
- ▶ Big SQL is a **massively parallel processing (MPP) SQL engine** that runs in Apache Hadoop to achieve vastly improved performance and SQL execution breadth over other SQL-on-Hadoop offerings.
- ▶ Big SQL 3.0 provides the following support over Hive 0.13:
  - ▶ More comprehensive SQL support (see below for details)
  - ▶ Federated queries
  - ▶ **Statistics-driven optimization** and **query planning**

Check the reference below for more detail feature comparison with Hive.

# Which one to choose?

- ▶ Each has its advantage and disadvantages. You have to find out what is best for your application.
- ▶ Things to consider
  - ▶ Data Access Pattern
  - ▶ Volume
  - ▶ Cost
  - ▶ performance.
- ▶ Following is an article I found online
  - ▶ <http://blog.markedup.com/2013/02/cassandra-hive-and-hadoop-how-we-picked-our-analytics-stack/>

# Hive Cheat Sheet for commands

## ▶ Retrieving Information (General)

▶ `SELECT from_columns FROM table WHERE conditions;`

▶ **Retrieving All Values** `SELECT * FROM table;`

▶ **Retrieving Some Values** `SELECT * FROM table WHERE rec_name = "value";`

## ▶ Retrieving With Multiple Criteria

▶ `SELECT * FROM TABLE WHERE rec1 = "value1" AND rec2 = "value2";`

▶ **Retrieving Specific Columns** `SELECT column_name FROM table;`

▶ **Retrieving Unique Output** `SELECT DISTINCT column_name FROM table;`

▶ **Sorting** `SELECT col1, col2 FROM table ORDER BY col2;`

▶ **Sorting Reverse** `SELECT col1, col2 FROM table ORDER BY col2 DESC;`

▶ **Counting Rows** `SELECT COUNT(*) FROM table;`

▶ **Grouping With Counting** `SELECT owner, COUNT(*) FROM table GROUP BY owner;`

▶ **Maximum Value** `SELECT MAX(col_name) AS label FROM table;`

# References

1. **Programming Hive** -By [Edward Capriolo, Dean Wampler, Jason Rutherglen](#) (O'Reilly Media)
  - ▶ Edward Capriolo, Dean Wampler & Jason Rutherglen
2. **Bigdatauniversity :**
  1. “Accessing Hadoop Data Using Hive”
  2. “SQL Access for Hadoop”
3. **Apache Hive tutorial** <https://cwiki.apache.org/confluence/display/Hive/Tutorial>
4. <http://hortonworks.com/hadoop-tutorial/how-to-process-data-with-apache-hive/>
5. <https://developer.ibm.com/hadoop/docs/getting-started/faqs/>



# THANK YOU!

*“Big data is the new Natural resource”* – Rometty, CEO IBM

ARE YOU READY to use this new resource?

- ▶ Let me know if you want to volunteer for any **presentation** on any topic in future meetup.
- ▶ Any suggestion for future **venue** and **topics**?

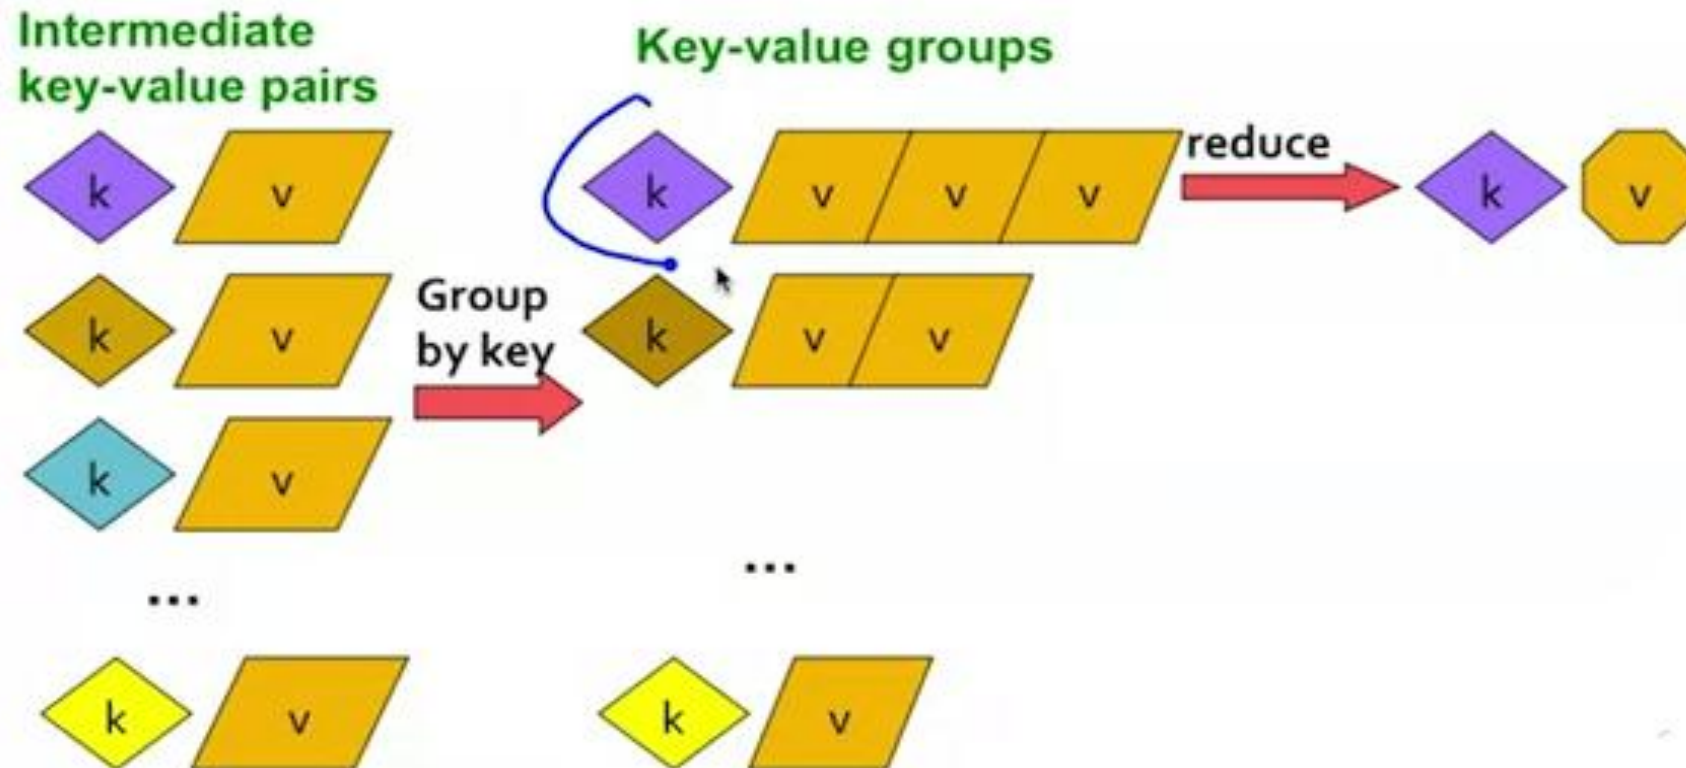
# Backup Slide

## ► Map - Reduce

- Store **data redundantly** on multiple nodes for persistence and availability
- Move **computation close to data** to minimize data movement
- **Simple programming model** to hide the complexity of parallel computing in clusters of computer

# Backup slide -Map Reduce Basics Revisited

► 32



# Code for word Count

```
CREATE TABLE docs ( line STRING);

LOAD DATA INPATH 'docs' OVERWRITE INTO TABLE docs;

CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
  ( SELECT explode(split(line, '\s')) AS word FROM docs) w
GROUP by word
ORDER by word;
```